



Steal the model and read a stored embedding



Model extraction: an inference API is a model you can copy

UC1: Unbounded Query Extraction via /chat → Behavioural Clone



Key insight: No control plane between attacker and model. Every call teaches the attacker the model's behaviour.

Attacker-controlled

Target / System

Passive capture

Decision logic



OWASP LLM10
Unbounded Consumption
(no rate limit, no metering → extraction at scale)




MITRE ATLAS
AML.T0024
Exfiltration via ML Inference API



STRIDE-for-LLM
Information Disclosure (+ DoS lever)
(confidentiality + availability risks)

Membership inference: confidence already leaks what's in the store



MIDDLE RUNG: nothing reconstructed yet — this only answers WHETHER a text is indexed




MITRE ATLAS AML.T0024
Exfiltration via ML Inference API (inference-time)



STRIDE-for-LLM: Information Disclosure

Embedding inversion: a stored vector is not a one-way hash



MIDDLE RUNG: nothing reconstructed yet — this only answers WHETHER a text is indexed




MITRE ATLAS AML.T0024
Exfiltration via ML Inference API (inference-time)




STRIDE-for-LLM: Information Disclosure


No tenant boundary: the tenant is known, passed, and thrown away




The demo: an outsider's toolkit against two doors




The acceptance run: one line that turns three attacks into a regression




Model extraction, at client scale: a clone copies the bugs, not just the answers



Embedding inversion, at client scale: the same script shows the leak and its fix



Membership inference at client scale: a leaked store answers "is this text ours?"



What we walk away with: two files, one line, and nothing fixed

