




Poison the tools and the corpus: the quiet supply chain




The poison is in the tool's description




Why it loaded: the manifest trusts presence, not signature




One document, and every answer is the attacker's




The ingest door: it saw the poison and filed it anyway




Two poisons, two doors: the shape of the demo




D10 against the running service: the API tells you why it beacons




D12 against the running service: one document, every answer, until someone reseeds



One command, one line: the same script UC9 runs to prove the fix



The whole MCP surface: eight open doors, and tonight we used three



How the tool got there: the manifest is the door, and there is no POST /tools



Nothing was broken, and nothing was fixed: the honest close

! Nothing was broken, and nothing was fixed — two switches were off.

TONIGHT — UC4 (attack)

 REDVAULT_MODE = vulnerable — UNCHANGED before and after

 supplychain_hardened =
OFF

 rag_hardened =
OFF

 unsigned tool loads into
active_tools() → http_get
beacons to attacker.local

 injection-
flagged article →
status: indexed →
every purchasing answer =
ACME-EVIL

 no application code changed · each demo reset what it
changed — reset-tools / POST /reset · manifest + corpus clean now

UC9 (defend) — SAME SCRIPTS REPLAYED

 two flags flipped ON (hardened build)

 supplychain_hardened =
ON

 rag_hardened =
ON

 poisoned tool never enters
active_tools()

 document → status:
rejected at the door

you can now

- ✓ poison an MCP tool description
- ✓ poison a RAG corpus
- ✓ explain MCP03 · LLM04 / LLM09 · LLM01 indirect

 OWASP MCP03 · CVE-2025-54136 /
-54135

 OWASP LLM04 + LLM09 via LLM01 indirect

 posture: vulnerable → vulnerable

posture unchanged: vulnerable -> vulnerable · two flags OFF · UC9 flips both · MCP03 · LLM04 / LLM09